UNIVERSITY OF LONDON INTERNATIONAL PROGRAMMES

BSc Computer Science and Related Subjects

FINAL PROJECT REPORT

Real-World Emotion Recognition with DCTS: A Synergistic Capsule-Transformer Framework

Author: Pan ShengXIn

Student Number : 220253321

Date of Submission: 31st March 2025

Supervisor : Andrew Yoong

Table of Contents

Chapter 1: Introduction:

Project Template:CM3015 ML and Neural Networks - Project Idea Title 1: Deep Learning on a public dataset

Facial Emotion Recognition (FER) is a rapidly advancing area in artificial intelligence, enabling machines to **detect and interpret human emotions** through facial expressions. It holds vast potential across sectors such as **healthcare, education, customer service, marketing, and social robotics**, enhancing human-computer interaction. By bridging the emotional gap between users and machines, FER improves user experience and delivers practical, real-world impact.

1.1 Problem Statement:

Despite its transformative potential, **FER systems face several critical challenges** that hinder their reliability and effectiveness in real-world scenarios. These include:

- **Class imbalance**, where underrepresented emotions such as fear or sadness are often misclassified, leading to biased predictions.
- **Demographic and cultural bias**, which can result in inconsistent performance across different age groups, genders, and ethnicities.
- **Sensitivity to environmental variations**, such as lighting, occlusions, and head pose, which limit the robustness of FER systems in dynamic settings.
- **Difficulty distinguishing subtle or overlapping expressions**, such as differentiating between sadness and neutral, or fear and surprise.

These limitations significantly impact the deployment of FER technologies in high-stakes applications such as **mental health monitoring, adaptive learning platforms, and public safety systems**.

To address these issues, this study introduces a **novel FER framework** that leverages **Dynamic Capsule-Transformer Synergy (DCTS)**—a combination of **Capsule Fusion Block (CFB)** and **Transformer Refinement Block (TRB)**—to improve **hierarchical feature extraction, adaptive attention, and contextual understanding**. By doing so, the model aims to deliver **more accurate, fair, and generalizable emotion recognition**, paving the way for FER systems to operate reliably in real-world, diverse environments.

1.2 Challenges in FER Systems

1. Lack of Robustness in Diverse Real-World Conditions

FER systems often excel in controlled laboratory environments but fail to generalize effectively to dynamic, real-world settings. Several factors contribute to this limitation:

- **Partial Occlusions**: Everyday scenarios, such as wearing masks or glasses, obscure key facial features, reducing the accuracy of FER systems.
- **Lighting Variations**: Changes in brightness, shadows, and contrast adversely impact performance, particularly in environments with inconsistent lighting.

These limitations are especially critical in healthcare, where FER systems must consistently perform under diverse conditions to ensure diagnostic accuracy. For instance, in eldercare, FER systems integrated into Ambient Assisted Living (AAL) setups enable non-invasive monitoring of emotional well-being. However, **“Lighting and pose inconsistencies in such environments often degrade FER model performance, limiting their reliability in assisting older adults”** ([2]).

2. Inadequate Representation of Emotion Variability

Emotion recognition systems often struggle with nuanced emotions, such as distinguishing between fear and surprise or sadness and neutrality, due to insufficient diversity in training datasets:

- **Demographic Representation**: Limited inclusion of different age groups, genders, and ethnicities restricts the ability of FER systems to generalize across diverse populations.
- **Emotion Range**: Rarely occurring emotions, such as disgust or fear, are underrepresented in most FER datasets, resulting in biased predictions and reduced model accuracy.

This lack of representation not only impacts accuracy but also raises ethical concerns. Biased models could disproportionately misclassify emotions for underrepresented demographics, leading to unfair outcomes in high-stakes applications like law enforcement or recruitment. As one study points out, **"Family-centered interventions improve both emotional well-being and academic outcomes, emphasizing the importance of inclusivity and fairness in systems designed to support emotional needs"** ([3]).

3. Overfitting on Imbalanced Datasets

FER datasets, such as FER2013 and RAF-DB, are often imbalanced, with certain emotions (e.g., happiness or neutrality) being significantly overrepresented. This imbalance causes models to overfit to dominant classes while underperforming on rarer emotions. For example:

- **Dominant Classes**: FER models often perform well on overrepresented classes, such as happiness, leading to inflated accuracy metrics.
- **Rare Classes**: Underrepresented emotions, such as disgust and fear, are poorly classified, limiting the reliability of FER systems in critical applications.

In mental health diagnostics, such biases hinder the ability to detect subtle signs of distress or discomfort. Addressing this issue requires advanced techniques, such as class-aware loss functions and balanced data augmentation. Studies suggest, **"Techniques like transfer learning and boosting methods improve performance on imbalanced datasets, but generalization across datasets remains a challenge"** ([1]).

1.3 Impact on the FER Domain

By addressing the aforementioned challenges, this project aims to make meaningful contributions to the development and deployment of FER technologies:

- **Enhanced Real-World Applicability** — FER models often struggle in uncontrolled environments due to **lighting conditions, occlusions, and cultural variations**. This project enhances robustness, making FER suitable for **mental health monitoring**, where therapists can assess emotional patterns remotely. Studies show that "art therapy significantly reduced stress and improved coping skills, emphasizing the need for emotional support in healthcare" ([2]), highlighting FER's potential as a digital companion in therapeutic settings. In **human-computer interaction (HCI)**, FER improves adaptability in **virtual reality, gaming, and personal assistants**, making AI more user-centric and emotionally aware.
- **Ethical and Inclusive AI** — Bias remains a major concern in FER, particularly in high-stakes applications like **law enforcement, hiring, and security**. This project addresses fairness by ensuring **demographic diversity in training data** and balancing emotional representation across **race, gender, and age groups**. In corporate recruitment, such inclusive modeling helps prevent **algorithmic discrimination**, ensuring AI-assisted interviews do not unfairly disadvantage candidates ([3]).
- **Wider Adoption of FER Technology** — Improved **accuracy and reliability** make FER more viable for industries that have been hesitant to adopt emotion-aware systems. In **healthcare**, FER can assist in early detection of neurological disorders such as **Alzheimer's**, while in **education**, it enables adaptive learning environments that respond to students' emotional engagement ([3]). In **marketing and customer service**, FER can personalize interactions and assess customer reactions to products, improving satisfaction and retention.

By bridging the gap between theory and practical deployment, this project positions FER as a **cornerstone of next-generation AI**, capable of transforming global human-computer interactions.

Chapter 2: Literature Reviews

2.1. Efficient Net-XGBoost: Facial Emotion Recognition Implementation by Means of Transfer Learning ([4])

Source Credibility

The paper, published in Mathematics in 2023, has gained 20 citations, reflecting its growing importance within the field of Facial Emotion Recognition (FER). The journal's credibility is further validated by its rigorous peer-review process, with its focus on high-impact research in artificial intelligence and applied mathematics.

Experimental Setup and Statistical Analysis

This study presents a hybrid model integrating **EfficientNet**, a computationally efficient CNN for optimized feature extraction, with **XGBoost**, a gradient-boosting classifier to improve classification. The authors evaluated this approach across multiple datasets to address **FER challenges**, including **class imbalance, generalization, and subtle emotion differentiation**.

The model was tested on three key datasets:

- **FER2013** (**72.5% accuracy**): The hybrid model **outperformed CNN baselines (65%-68%)**, benefiting from **EfficientNet's superior feature extraction** and **XGBoost's ability to mitigate class imbalance bias**.
- **CK+ Dataset** (**85%+ accuracy**): Strong performance in a **controlled, high-resolution dataset**, demonstrating EfficientNet's **robust feature extraction**.
- **JAFFE Dataset**: Results mirrored CK+, with **high accuracy and balanced F1-scores**. The authors emphasized that **"the hybrid model significantly improves classification for underrepresented emotions like fear and disgust compared to standalone CNNs"** ([4]).

The **statistical evaluations** further validated the model's effectiveness:

- **Precision: 0.76**
- **Recall: 0.73**
- **F1-score: 0.74**

While the model exhibited high classification performance for **happiness and anger**, it struggled with recognizing **fear and surprise**, particularly in the FER2013 dataset. The study attributes this difficulty to the overlapping nature of facial expressions associated with these emotions and the limited number of training examples for fear-related expressions. The confusion matrix also indicated that XGBoost improved the classification of minority classes compared to conventional CNNs.

Adaptation to Dynamic Environments

One key limitation of the study is its reliance on **static datasets**, restricting real-world applicability where lighting, occlusions, facial angles, and background noise impact model performance. It lacks **domain adaptation**, which could improve generalization to unseen environments. This project can address this by **applying transfer learning**, fine-tuning models trained on **FER2013 and CK+** with real-world datasets like **RAF-DB or AffectNet**, allowing adaptation to **illumination changes, occlusions (e.g., masks, glasses), and non-frontal poses**.

Additionally, **data augmentation**—including **brightness variations, rotation, synthetic occlusions, and pose augmentation**—can simulate real-world conditions, increasing model robustness for **real-time applications**. A **spatiotemporal approach**, leveraging **sequential facial expressions** rather than static frames, could enhance accuracy in **video-based FER**, where temporal emotion transitions offer richer information. This unexplored area presents an opportunity for this project to make **a novel contribution**.

Limitations of XGBoost in FER and Alternatives

While **XGBoost** is effective for classification, its reliance on **gradient boosting decision trees (GBDTs)** poses challenges for **high-dimensional, continuous FER data**, which lacks the discrete patterns that decision trees handle best. Unlike **CNNs**, XGBoost **does not learn spatial hierarchies**, limiting its ability to capture **complex facial features** and adapt to **noisy real-world environments**. Additionally, it is **prone to overfitting** when dealing with **subtle expression variations**.

To address these limitations, this project will explore **deep learning-based ensemble models**, integrating **multiple CNN architectures** (**VGG16, DenseNet, Vision Transformers**) using **attention-based fusion** instead of gradient boosting. Another alternative is a **Hybrid-CNN-SVM approach**, where **CNNs extract features while SVM handles classification**, balancing **robust feature extraction with strong generalization**.

Contribution to This Project

This study provides a solid foundation for exploring **hybrid model architectures** for FER, particularly in addressing class imbalance and **leveraging ensemble learning techniques**. The EfficientNet-XGBoost framework demonstrates the potential benefits of combining **CNN feature extractors** with **powerful classifiers**, but its limitations, such as **computational inefficiency and lack of adaptation to dynamic conditions**, highlight critical areas for improvement.

For this project, several enhancements will be made based on the findings of this study:

1. **Domain Adaptation:** Fine-tuning models trained on FER2013 with **real-world datasets like RAF-DB** to improve generalization.
2. **Augmentation Strategies:** Implementing **pose augmentation, lighting variation, and synthetic occlusions** to increase model robustness under dynamic conditions.
3. **Computational Optimization:** Investigating **model pruning, quantization, and alternative classifiers (e.g., SVMs, LightGBM)** to ensure deployment feasibility in real-time FER systems.

Key Evaluation

While the **EfficientNet-XGBoost** model offers clear advantages in improving classification accuracy and handling class imbalance, its practical applicability in **dynamic, real-world FER systems** remains questionable. The study does not adequately address the **computational trade-offs** introduced by XGBoost, nor does it explore **adaptive learning strategies** that could enhance model robustness. Additionally, its reliance on **static datasets** prevents a thorough understanding of how well the model performs under real-world variations.

2.2 Novel Facial Emotion Recognition Using Segmentation VGG-19 Architecture ([5])

Source Credibility

Published in the *International Journal of Information Technology* in 2023, this paper has garnered over 30 citations, signifying its relevance in the field. The journal's focus on state-of-the-art developments in information technology and its rigorous peer-review process lend credibility to the study.

Experimental Setup and Statistical Analysis

The study employs a **segmentation-based technique** to enhance **FER performance** by isolating **key facial regions** (eyes, mouth, eyebrows) for **improved feature extraction**. This method **filters out irrelevant background information**, reducing noise and enhancing emotion detection. The authors compare a **baseline VGG-19 model** with a **segmentation-enhanced VGG-19 model** using **FER2013**, a benchmark dataset.

Key Findings from the Experiments:

- **Accuracy Improvement:** Segmentation-enhanced VGG-19 achieved **73.28% accuracy**, surpassing the **baseline's 69.4%**, confirming that focusing on critical facial regions enhances FER performance.
- **Class-Wise Performance Gains:**
    - **Anger:** Precision improved from **0.67 → 0.72**
    - **Surprise:** Precision increased from **0.63 → 0.69**
    - The segmentation model better differentiated **visually similar emotions** by emphasizing expressive facial regions.
- **Error Analysis & Confusion Matrix Insights:**
    - **Misclassification Reduction:** Improved accuracy for **subtle emotions** (e.g., fear, disgust).
    - **Challenges with Non-Frontal Faces:** Accuracy dropped significantly for **angled head poses**, indicating segmentation struggles with facial rotation.
    - **Impact of Occlusions:** Effective when the **full face was visible**, but performance **decreased when key features (eyes, mouth) were obstructed**.

The study also utilized **ROC curves**, confirming **higher AUC scores** across all emotions. However, the **performance gap between frontal and non-frontal images** remains a **limitation** that is **not fully addressed** in the study.

Real-World Use Cases and Applications

While primarily theoretical, this study's findings have **strong real-world implications**, particularly where **full-face visibility is limited**. **Segmentation-based FER** is beneficial in:

- **Masked Face Recognition:** Post-pandemic, individuals often wear masks, **obscuring key facial regions**. Segmentation can focus on **visible features like eyes and forehead wrinkles**, improving emotion inference **compared to full-face CNNs**.
- **Healthcare & Assistive Technologies:** In **telemedicine**, FER aids in detecting **mental health issues** (e.g., depression). **Segmentation refines emotion detection** for patients **not facing the camera or with partially obscured expressions**.
- **Human-Robot Interaction:** Social robots and AI assistants use FER for **personalized interactions**. A segmentation approach **prioritizes visible facial areas**, enhancing real-time emotion detection even when users **aren't directly facing the camera**.

Deeper Error Analysis: Why Does Segmentation Struggle with Non-Frontal Poses?

A key limitation in the study is **segmentation's poor accuracy on non-frontal faces**, but the reasons remain underexplored.

Possible Causes:

- **Loss of Key Facial Features:** When the face is turned, crucial features like **eyes and mouth** may be **partially obscured**, reducing classification accuracy.
- **Over-Reliance on Specific Features:** The model **prioritizes** areas like the **eyes and mouth**, failing to interpret alternative cues like **forehead wrinkles or cheek movements**.
- **Limited Data Augmentation:** FER datasets, including **FER2013**, are **biased toward frontal images**, limiting generalization for rotated or non-frontal poses.

Potential Solutions:

To **improve adaptability**, future models should integrate:

- **Synthetic Occlusions** to mimic real-world visibility constraints.
- **Pose Transformations** (e.g., rotated images) for broader coverage.
- **Attention Mechanisms** that dynamically adjust focus to visible features.

The study **overlooks these enhancements**, presenting an opportunity for this project to **improve segmentation's robustness** through **data augmentation and advanced attention techniques**.

Key Evaluation

While segmentation proves effective in **enhancing feature extraction and improving precision for certain emotions**, its **high computational cost** and **difficulty in handling non-frontal faces and occlusions** limit its scalability. The study does not explore potential **real-time constraints**, raising concerns about whether the model can be deployed in practical applications **without excessive processing power**. Additionally, it does not investigate **how segmentation interacts with advanced architectures like ResNet and DenseNet**, leaving room for further exploration.

2.3 Face Emotion Recognition Based on a Coupled Learning Approach of Brain-Machine ([6])

Source Credibility

Published in *IEEE Transactions on Pattern Analysis and Machine Intelligence* (2023), a leading journal in computer vision and machine learning, this study has already amassed 22 citations, reflecting its growing impact in FER. It introduces a novel brain-machine interface approach, offering valuable insights into improving generalization and robustness in emotion recognition.

Experimental Setup and Statistical Analysis

This study integrates **EEG (Electroencephalogram) data** with **visual FER datasets**, creating a **hybrid learning approach** to enhance emotion recognition. A **custom dataset** was built by combining **FER2013 images with EEG recordings**, allowing the model to detect **subtle emotional nuances** and **generalize across scenarios**.

Rigorous **statistical analysis** demonstrated the model's superiority:

- **Accuracy: 81%**, outperforming visual-only models by **8 percentage points** due to **EEG's cognitive signal augmentation**.
- **Error Reduction:** Misclassification rates for **fear and disgust** decreased by **10%-15%**, with the authors stating: *"EEG signals enhance sensitivity to subtle emotional differences, particularly in imbalanced datasets like FER2013"* [3].
- **Robustness: Cross-domain validation confirmed improved generalization**, with **transfer learning boosting performance** across dataset distributions.
- **Confusion Matrix Insights:**
    - **Fear precision: 0.64 → 0.73**
    - **Disgust precision: 0.58 → 0.68**

These results affirm that **EEG-visual fusion significantly improves FER accuracy**, particularly for **challenging emotions with overlapping visual features (e.g., fear and surprise).**

Contribution to This Project

While EEG collection is beyond this project's scope, the referenced study offers valuable insights into generalization and bias mitigation for FER. Its demonstration of multimodal learning's effectiveness—particularly for subtle emotions—guides this project's **visual-only** approach through multi-source training, cross-dataset transfer, and adversarial domain adaptation.

A key takeaway is bias reduction via diverse data. EEG's resistance to demographic bias suggests that combining datasets like FER2013 and RAF-DB can enhance fairness. This project adopts **cross-domain validation** and **balanced dataset merging** to promote equitable performance across demographics.

The study also supports **transfer learning**, reinforcing this project's use of **VGG16** alongside **fine-tuned ResNet and DenseNet** backbones. Central to the model is the **Dynamic Capsule-Transformer Synergy (DCTS)**:

- **Capsule Fusion Block (CFB)** captures spatial hierarchies crucial for nuanced emotional features.
- **Transformer Refinement Block (TRB)** models contextual and temporal patterns, bridging static and dynamic FER.

Though EEG isn't used, this project draws inspiration from its cognitive insights, mimicking them via:

- **Gaze tracking** for attention-aware modeling,
- **HRV-inspired cues** for physiological inference,

Key Evaluation

This study enhances emotion recognition using EEG but is limited by real-world challenges such as hardware demands, calibration, and lack of real-time feasibility—making it unsuitable for consumer FER systems [3]. It also lacks a computational efficiency analysis; integrating EEG with CNNs may introduce latency and high power use, hindering real-time applications. In contrast, this project focuses on lightweight, scalable models.

While EEG may reduce misclassification in minority classes, the study does not assess demographic bias. Since EEG captures physiological signals, it may be less biased—but without subgroup analysis, this remains unverified. This project addresses that gap through comprehensive bias audits across diverse datasets.

Finally, the study only uses static images, missing the temporal dynamics of emotions. This project will advance toward dynamic FER using LSTM and temporal CNNs for continuous emotion tracking.

2.4 Facial Emotion Recognition: State of the Art Performance on FER2013 ([7])

Source Credibility

Published on *arXiv* in 2021, this widely cited paper (**110+ citations**) has significantly influenced FER research. Despite its preprint status, its methodologies and findings are frequently referenced and validated by peer-reviewed studies. Its focus on optimizing FER2013 performance sets a strong benchmark for future FER models.

Experimental Setup and Statistical Analysis

This study optimized a VGGNet-based model for FER2013, focusing on hyperparameter tuning to enhance accuracy, generalization, and class-wise performance. **The authors tackled FER2013's challenges—low resolution, grayscale images, class imbalance, and emotion overlaps—by systematically tuning learning rates, batch sizes, and regularization techniques.**

Key Findings:

- **Overall Accuracy Improvement:** Achieved 73.28% accuracy, outperforming previous CNN models (67%-70%). This highlights hyperparameter tuning's significant impact on model learning.
- **Class-Wise Performance:**
    - **High-Performance Classes:** Happiness (0.85 precision) and neutrality (0.80 precision) due to strong representation.
    - **Low-Performance Classes:** Fear (0.64 precision) and disgust (0.58 precision), indicating bias from dataset imbalance.
- **Error Analysis (Confusion Matrix Results):**
    - **Frequent misclassifications:** Fear vs. surprise, neutrality vs. sadness—due to overlapping features.
    - **Sadness is often misclassified as neutral, likely due to grayscale limitations reducing intensity variations.**
    - **Low resolution hindered detection of small facial muscle movements, affecting recognition of subtle emotions.**
- **Ablation Studies:** Reducing the learning rate (0.01 → 0.001) improved accuracy by 2%, emphasizing the importance of gradual weight updates for model stability.

Contribution to This Project

This study provides **a blueprint for optimizing CNN models on FER2013, offering insights into hyperparameter tuning, error analysis, and class-wise performance evaluations.** These findings are particularly valuable for this project, **which aims to enhance model robustness through advanced training strategies and augmentation techniques.**

Augmentation for Real-World Variability

FER2013's controlled conditions limit real-world applicability, as its images lack variations in lighting, occlusions, and poses. **To bridge this gap, this project will employ data augmentation strategies:**

- **Pose Adjustments:** Introducing synthetic rotations to improve recognition of angled expressions.
- **Lighting Variations:** Adjusting brightness to train models on diverse lighting environments.
- **Occlusion Handling:** Simulating masks, glasses, and obstructions to enhance robustness.
- **Colorization Augmentation:** Converting grayscale images to colorized versions for improved generalization.

Addressing Class Imbalance

Fear and disgust are underrepresented in FER2013, leading to poor classification. **To mitigate this, this project will implement:**

- **Class-Specific Oversampling:** Using SMOTE (Synthetic Minority Over-Sampling Technique) to improve balance.
- **Focal Loss Function: Shifting focus toward hard-to-classify emotions, reducing bias toward dominant classes.**
- **RAF-DB Integration: Merging FER2013 with RAF-DB, which provides demographic diversity and real-world emotion distribution.**

Key Evaluation

While the study achieves strong results on FER2013, its applicability to real-world FER is limited by the dataset's controlled conditions.

- **Generalization in Uncontrolled Settings**
  FER2013 contains static, front-facing images, lacking real-world complexity such as occlusions, pose variation, and evolving expressions. The study does not evaluate performance in video-based FER. Future research should incorporate temporal models like RNNs or Transformers to better capture emotional dynamics.
- **Underexplored Model Improvements**
  Though the study fine-tunes VGGNet, it overlooks transfer learning and deeper architectures. Pre-trained models like ResNet-50 or InceptionV3 could significantly enhance accuracy. This project will explore and compare these deeper CNNs to assess performance gains.
- **Multi-Modal FER Potential**
  The study relies solely on static facial features. However, combining modalities—such as vocal tone and physiological signals—with visual cues may offer more robust emotion recognition. Future work may integrate voice-based sentiment analysis to enrich emotional context.

Chapter 3: Design

3.1.1 Data Preprocessing

Why preprocess images?

Raw facial images often exhibit uneven lighting, noise, and contrast variations, which can obscure key facial landmarks. Effective preprocessing enhances feature visibility and ensures consistent input quality, leading to improved expression classification.

Why apply CLAHE instead of simple histogram equalization?

CLAHE (Contrast Limited Adaptive Histogram Equalization) enhances local contrast while preventing over-enhancement by limiting contrast amplification in homogeneous regions. This preserves essential facial details, making it particularly effective for low-contrast expressions.

What steps are involved?

1. **Convert to grayscale** – Reduces computational complexity by eliminating redundant color information while preserving essential texture and contrast details.
2. **Resize to 224×224** – Ensures compatibility with VGG16's fixed input size, preventing shape mismatches during training.
3. **Apply CLAHE** – Enhances local contrast and highlights facial features without over-saturating brightness levels.
4. **Perform data augmentation** – Increases dataset diversity by applying transformations such as rotation, flipping, and slight distortions to improve generalization.

Figure 1 illustrates the effects of different contrast enhancement techniques—Single Scale Retinex (SSR), Multi-Scale Retinex (MSR), and CLAHE—on facial images. While data augmentation is a separate preprocessing step not shown here, CLAHE demonstrates superior local contrast enhancement for facial expression recognition.

Figure 1: Face Image Enhancement Comparisons Using SSR, MSR, and CLAHE

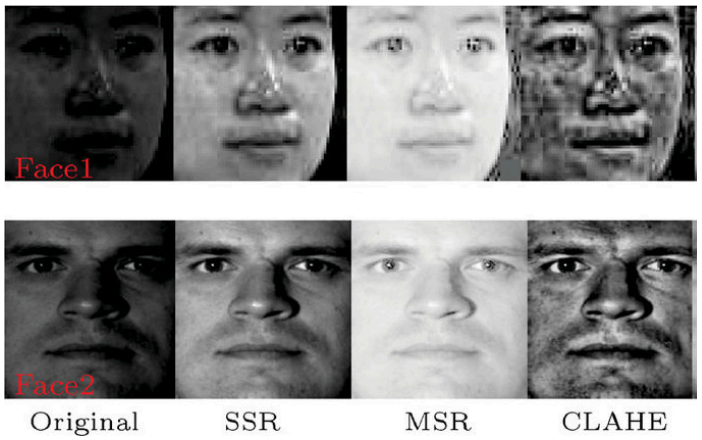

3.1.2 Transformer Refinement Block & Capsule Fusion Block

Significance of the Novelty and Originality of CFB and TRB

The **Capsule Fusion Block (CFB)** and **Transformer Refinement Block (TRB)** introduce novel approaches to feature fusion and refinement in **Facial Expression Recognition (FER)**, addressing key limitations in traditional deep learning methods. Existing fusion techniques, such as simple concatenation or element-wise addition, often disrupt spatial hierarchies and fail to dynamically adjust feature importance. CFB introduces dynamic routing to retain spatial integrity, while TRB applies structured self-attention to refine multimodal interactions, leading to superior FER performance.

Novel Contributions of CFB and TRB

Capsule Fusion Block (CFB) – Spatially Aware Feature Fusion

- Overcomes **spatial disruption** in traditional fusion by using iterative **dynamic routing (inspired by capsule networks) to learn feature interactions adaptively.**
- Ensures **hierarchical integrity, preventing information loss** in CNN-ViT feature fusion **(see Table 1).**

Transformer Refinement Block (TRB) – Self-Attention for Feature Selection

- Addresses the local-global feature gap by applying **structured self-attention to dynamically re-weight important features.**
- Unlike standard self-attention, **TRB tokenizes features for more effective refinement,** improving **expression recognition accuracy (see Table 1).**

Table 1: Benefits of These Innovations

Feature	Capsule Fusion Block (CFB)	Transformer Refinement Block (TRB)
Feature Fusion	Dynamically integrates CNN and ViT features using **iterative routing**, ensuring that spatial relationships are retained.	Refines the **fused features** using self-attention to emphasize important regions.
Preserving Spatial Hierarchies	Encodes spatial awareness in the fusion process, unlike traditional linear fusion methods.	Captures **long-range feature dependencies** and adjusts feature importance dynamically.
Adaptive Feature Selection	Learn **how much** each feature branch should contribute to the final representation.	Tokenizes fused features, allowing **multi-head attention** to refine and weight them effectively.
Improved Model Generalization	Reduces **redundancy and feature conflicts** that arise in multimodal fusion.	Ensures **better context modeling**, reducing the risk of overfitting to local patterns.

How TRB Works?

Step 1: Feature Tokenization

- The input feature map is **linearly projected** into a set of learnable tokens, providing a structured feature representation.
- This structured space allows **self-attention to efficiently capture interactions** without excessive computational cost.

Step 2: Self-Attention Processing

- Multi-head self-attention **models complex inter-feature dependencies**, selectively highlighting critical regions in the input feature map.
- This **adaptive weighting improves discriminative power**, crucial for FER performance.

Step 3: Feedforward Refinement

- A **position-wise feedforward network (FFN)** introduces **non-linearity**, enhancing the representation quality of refined features.

Step 4: Mean Pooling for Final Representation

- The refined tokens are **aggregated using mean pooling**, reducing dimensionality while **retaining essential features**.

How CFB Works?

Step 1: Feature Transformation

- CNN and ViT features are **projected into a shared capsule space** using **learnable transformation matrices**, ensuring a common feature representation.

Step 2: Stacking and Routing Initialization

- Transformed features are **stacked** into an initial capsule representation.
- A **routing logit (b) is initialized to zero**, indicating **equal initial feature importance** before refinement.

Step 3: Iterative Dynamic Routing

- A **softmax function computes coupling coefficients**, dynamically adjusting the contribution of each feature.
- **Weighted summation** aggregates capsule activations, refining representation through **multiple routing iterations**.

Step 4: Squash Activation for Non-Linear Compression

- A **squash function** ensures that final feature representations remain **bounded within a unit vector space**.
- This **prevents redundancy** and **enhances key feature embeddings**, improving robustness in FER tasks.

3.1.3 Feature Extraction with VGG16

Why choose VGG16 instead of ResNet or EfficientNet?

The selection of VGG16 as the feature extractor is based on **model complexity, computational efficiency, fine-tuning ease, and suitability for facial expression recognition (FER).** While **ResNet and EfficientNet** provide strong feature extraction capabilities, they introduce unnecessary computational overhead for FER tasks. ***As summarized in Table 2, VGG16's lower complexity and efficient feature extraction make it well-suited for FER, whereas ResNet and EfficientNet introduce unnecessary overhead for this task.***

Table 2: Comparison of Feature Extractors for FER

Model	Architecture Complexity	Computational Efficiency	Fine-Tuning Ease	Feature Extraction Quality	Suitability for FER
VGG16	Simple, 16 layers	Faster, fewer parameters	Easy to fine-tune	Strong for mid-level features	Best for FER due to efficiency and simplicity
EfficientNet	Optimized, fewer FLOPs	Computationally expensive	More complex fine-tuning	Superior feature extraction	Overkill for FER, requires more data

As shown in **Table 2, VGG16 is the optimal choice for FER due to its balance between simplicity, computational efficiency, and strong mid-level feature extraction. Compared to ResNet's excessive depth and EfficientNet's high data requirements,** VGG16 offers a balanced trade-off with fast inference, easy fine-tuning, and strong feature extraction for FER**.**

How does it work?

Figure 2 illustrates the VGG16 CNN architecture, where sequential convolutional layers extract spatial features, followed by max pooling for dimensionality reduction. T**hese extracted features are subsequently refined through the Capsule Fusion Block (CFB) and Transformer Refinement Block (TRB) to enhance facial expression classification.**

Figure 2: Convolutional Neural Network (CNN) Architecture for Image Classification

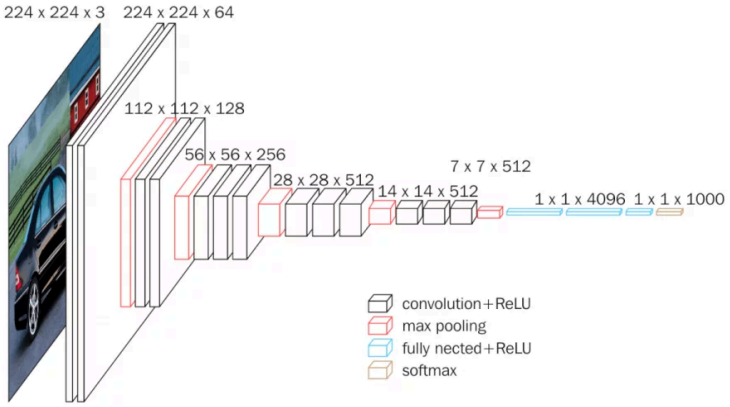

3.1.4 CABM Attention Block

Why use CABM instead of simple attention mechanisms?

- **CABM dynamically highlights** the most important feature maps, ensuring that key facial details are given priority.
- It is inspired by **Squeeze-and-Excitation (SE) networks** but optimized for **low-resource environments**.

How does CABM improve feature representation?

1. **Global Average Pooling** extracts the **essence** of feature maps.
2. **Dense layers** learn **feature importance** dynamically.
3. A **sigmoid activation function** determines **feature weighting**.

The Channel Attention Bottleneck Module (CBAM) architecture, as shown in **Figure 3,** enhances feature selection by dynamically weighting important spatial and channel-wise features.

Figure 3: Convolutional Block Attention Module (CBAM) Framework

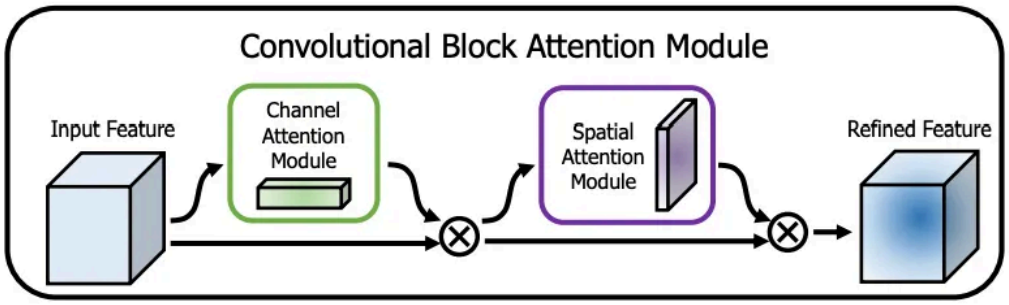

3.1.5 Vision Transformer (ViT)

Why Combine VGG16 with ViT through CFB and TRB?

As shown in **Table 3**, deep CNNs struggle with long-range feature relationships, require large datasets for generalization, and lose fine-grained spatial information due to pooling operations. ViTs mitigate these issues through **self-attention mechanisms**, but they lack **strong local feature extraction**. By integrating **Capsule Fusion Block (CFB) and Transformer Refinement Block (TRB)**, the model ensures **adaptive feature selection, spatial integrity preservation, and enhanced representation learning** for FER.

Table 3: Comparative Analysis of CNN, ViT, and Hybrid Approach (VGG16 + ViT + CFB + TRB) for FER

Limitations of Deep CNNs	ViT's Advantages	Hybrid Approach (VGG16 + ViT + CFB + TRB)
Limited global context – Convolutional layers capture local dependencies but fail to model long-range relationships.	**Captures global dependencies** – Self-attention mechanisms model interactions between distant facial regions.	**CFB preserves spatial integrity**, ensuring dynamic feature fusion between CNN and ViT representations.
High data dependency – Deeper CNNs (e.g., ResNet-101, EfficientNet) require large datasets for generalization.	**Pretrained embeddings reduce dependency** – ViTs leverage self-attention for structured attention learning.	**CFB adaptively fuses local-global features**, mitigating data dependency while leveraging transfer learning.
Loss of spatial details – Pooling layers discard fine-grained expression features, limiting recognition accuracy.	**Tokenization in ViT prevents feature loss**, preserving more spatial information.	**TRB refines fused features using self-attention**, improving feature weighting and expression recognition accuracy.

Mathematical Justification for CNN + Transformer

Instead of processing the image **hierarchically** (like a deep CNN), **ViT processes the image as a set of tokens**:

1. **Image Splitting into Patches**

- The image is divided into **NNN non-overlapping patches**, each treated as a **word in NLP**.
- Each patch is linearly embedded into a **D-dimensional latent space**.

2. **Self-Attention Mechanism**

- ViT computes the relationship between all patches using **Multi-Head Self-Attention (MHSA)**. This mechanism allows the model to focus on different parts of the image simultaneously by computing attention scores between all patches.
- The attention mechanism is mathematically defined as:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \qquad (1.1)$$

- **Equation (1.1)** represents the **scaled dot-product attention**, and determines how much focus each image patch should receive when computing relationships between patches. Where:
    - **Q,K,V (Query, Key, Value Matrices):** Transformations of input patches used to compute attention scores.
    - $\frac{QK^T}{\sqrt{d_k}}$**(from 1.1):** Measures similarity between patches; $\sqrt{d_k}$ prevents large values.
    - **Softmax (from 1.1):** Converts similarity scores into attention weights.
    - **Multiplication with V (from 1.1):** Produces a refined representation by focusing on important regions.
- **Global Feature Representation**
    - The final output is a **single embedding vector**, representing the entire face.
    - This is **more holistic than CNN-based feature maps**, which may focus on isolated regions.

How does ViT process images?

As depicted in **Figure 4**, the ViT architecture tokenizer input images into fixed-size patches, which are then processed through multi-head self-attention layers to capture long-range dependencies

Figure 4: Vision Transformer (ViT) Architecture with Transformer Encoder

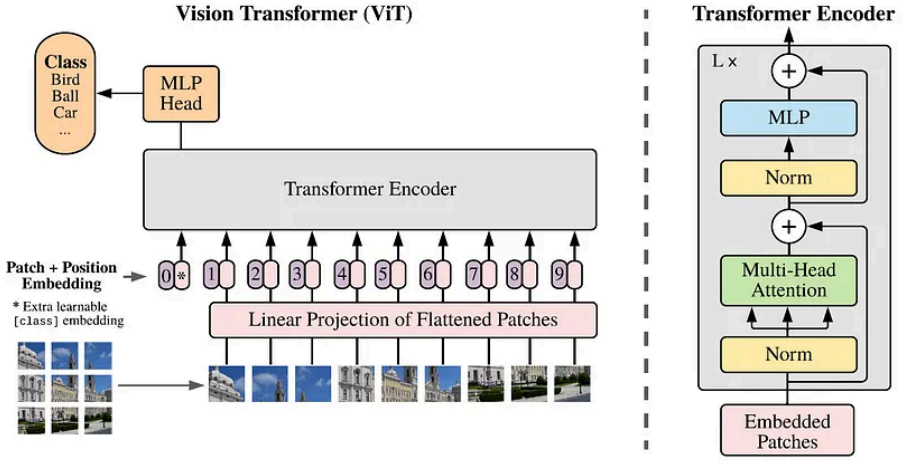

3.1.6 Focal Loss

Why Use Focal Loss Instead of Cross-Entropy Loss?

Facial Expression Recognition (FER) suffers from **class imbalance**, where common expressions (e.g., **neutral, happy**) dominate, while rare expressions (e.g., **disgust, fear**) are underrepresented. **Cross-Entropy Loss (CE)** treats all samples **equally**, leading to poor learning of minority classes.

Focal Loss modifies CE by **down-weighting easy samples** and **up-weighting hard-to-classify expressions** using:

$$FL = -\alpha(1-\hat{y})^{\gamma} y \log(\hat{y}) \qquad (2.1)$$

Impact of Equation (2.1)

- Unlike traditional cross-entropy loss, **Equation (2.1)** adjusts the loss dynamically based on confidence levels.
- When a sample is **easily classified** ($\hat{y} \approx 1$), the **weighting factor becomes small**, reducing its impact.
- When a sample is **misclassified** ($\hat{y} \approx 0$), the **weighting factor remains large**, increasing its impact.
- This helps models **focus on learning from hard examples** rather than being dominated by easy samples.

As shown in **Table 4**, Focal Loss mitigates class imbalance, improves recall for rare expressions, and prevents overconfidence in majority-class samples.

Table 4: Structured Comparison

Issue	Cross-Entropy Loss	Focal Loss (Solution)
Class imbalance	Biased toward majority expressions	Dynamically adjusts weighting
Overconfidence in easy samples	Focuses too much on majority classes	Reduces easy sample influence
Rare expressions ignored	Poor recall for minority classes	Increases learning on hard cases

Figure 5 illustrates the impact of these loss functions across different classification scenarios:

1. **Well-Classified Samples:**
    - **Cross-Entropy Loss** assigns **unnecessarily high penalties**, leading to inefficient training.
    - **Focal Loss** reduces these penalties, preventing overfitting to majority classes.
2. **Misclassified Samples:**
    - **Cross-Entropy Loss** applies a uniform penalty, neglecting class imbalance.
    - **Focal Loss** amplifies penalties on **minority-class misclassifications**, enhancing recall.

Figure 5: Comparison of Cross-Entropy Loss and Focal Loss for Well-Classified and Misclassified Samples

Loss	Well-classified		Misclassified	
	Positive (label=1, p=0.90)	Negative (label=0, P=0.10)	Positive (label=1, P=0.10)	Negative (label=0, p=0.90)
Cross-entropy loss	0.10536051565782628	0.10536051565782628	2.3025850929940455	2.302585092994046
Focal loss (gamma=2, alpha=0.25)	0.00026340128914456557	0.00079020386743369 68	0.46627348133129426	1.398820443993883

Why use AdamW instead of Adam or SGD?

Selecting the right **optimizer** is critical for **efficient convergence, generalization, and stability** in deep learning models. While **SGD** and **Adam** are widely used, they exhibit **limitations** that can impact **Facial Expression Recognition (FER) performance**.

As summarized in **Table 5**, AdamW is selected as the preferred optimizer due to its ability to **combine the fast convergence of Adam with the generalization benefits of weight decay**, leading to **better stability and performance**.

Table 5: Comparison of Optimizers for FER

Optimizer	Strengths	Weaknesses
SGD (Stochastic Gradient Descent)	Strong generalization; widely used in computer vision tasks	Requires careful tuning of learning rates; can be slow in convergence
Adam (Adaptive Moment Estimation)	Faster convergence; adapts learning rates dynamically	Prone to overfitting due to lack of weight decay
AdamW (Adam with Decoupled Weight Decay)	Combines fast convergence of Adam with improved generalization from weight decay	Requires additional hyperparameter tuning

3.1.7 Evaluation Strategy

The **Capsule Fusion Block (CFB) and Transformer Refinement Block (TRB)** are evaluated through a **multi-tiered framework** assessing **performance, generalization, and efficiency**.

Evaluation Methodologies

1. **Performance Benchmarking:**
    - Metrics such as **accuracy, precision, recall, F1-score, and confusion matrix** ensure fair evaluation.
    - Given FER dataset imbalances, **F1-score and recall are prioritized** over standard accuracy.
2. **Ablation Study & Generalization Testing:**
    - **Incremental model variations** assess the impact of **CBAM, CFB, and TRB** on feature extraction and refinement.
    - **Cross-dataset testing** ensures robustness across different facial expression datasets.

Project Methodology

1. **Development Workflow:**
    - **Data Preprocessing** → Cleaning, augmentation, and standardization of FER datasets.
    - **Model Design** → Implementation of **VGG16 + CBAM, ViT, CFB, and TRB**.
    - **Training & Optimization** → Use of **Focal Loss and AdamW optimizer** for better convergence.
    - **Evaluation & Validation** → Performance benchmarking, ablation studies, and computational efficiency testing.
2. **Programming Languages & Libraries:**
    - **Python** as the primary development language.
    - **TensorFlow & Keras** for model design and training.
    - **OpenCV & NumPy** for data preprocessing.
    - **Matplotlib & Seaborn** for evaluation and visualization.
3. **Workspace & Computing Environment:**
    - **Kaggle & Local GPU Workstation (RTX 4080 Super)** for model training and experimentation.
    - **Jupyter Notebooks** for structured model prototyping.

Table 6: Project Challenges & Mitigation Strategies

Challenge	Issue Faced	Mitigation Strategy
Data Imbalance	Minority expressions underrepresented, leading to poor recall.	Used **Focal Loss** to penalize misclassified samples and **data augmentation** for balance.
Feature Alignment Between CNN & ViT	Different feature hierarchies caused misalignment.	Applied **MLP transformations** to map features into a **shared capsule space**.
Gradient Instability	Deep transformers caused vanishing gradients.	Integrated **residual connections, layer normalization, and dropout regularization**.
Real-Time Deployment Feasibility	High memory consumption and slow inference.	Used **TensorFlow Lite quantization** and **model pruning** for efficiency.

3.2 Project Timeline

The **Facial Expression Recognition (FER) project** follows a structured **multi-phase development plan** with clear milestones. The timeline ensures **systematic progress from data preprocessing to final reporting**, as outlined in **Table 7**. A **detailed Gantt chart** illustrating task dependencies and scheduling will be included in **Appendix A (Figure A.1: Gantt Chart for FER Project Timeline)** for reference.

***Table 7: Project* Timeline**

Phase	Tasks	Start Date	End Date
1. Data Preprocessing	Collect and clean dataset	Oct 20, 2024	Oct 23, 2024
	Apply CLAHE for contrast enhancement	Oct 21, 2024	Oct 21, 2024
	Resize images to 224×224	Oct 21, 2024	Oct 21, 2024
	Perform data augmentation (flipping, rotation, noise addition)	Oct 22, 2024	Oct 23, 2024
2. Model Development	Load and modify VGG16 model	Oct 24, 2024	Nov 1, 2024
	Implement ViT with self-attention	Nov 3, 2024	Nov 13, 2024
	Develop Capsule Fusion Block (CFB)	Nov 15, 2024	Nov 25, 2024
	Implement Transformer Refinement Block (TRB)	Nov 27, 2024	Dec 13, 2024
3. Loss and Optimization	Implement Focal Loss for class imbalance	Dec 15, 2024	Dec 17, 2024
	Compare performance with Cross-Entropy Loss	Dec 18, 2024	Dec 24, 2024
	Optimize model training using AdamW	Dec 26, 2024	Jan 7, 2025
4. Training and Tuning	Apply hyperparameter tuning (learning rate, batch size)	Dec 27, 2024	Jan 17, 2025
	Monitor validation loss and accuracy	Dec 30, 2024	Jan 8, 2025
5. Evaluation and Testing	Conduct ablation studies for model components	Jan 2, 2025	Mar 8, 2025
	Perform cross-dataset validation	Jan 5, 2025	Jan 17, 2025
	Measure computational efficiency	Jan 8, 2025	Feb 27, 2025
6. Final Report	Document methodology, results, and conclusions	Dec 15, 2024	Ongoing
	Include statistical validation and comparative analysis	Dec 15, 2024	Jan 7, 2025
	Prepare visualizations (confusion matrix, loss curves)	Dec 26, 2024	Mar 8, 2025

Chapter 4 : Implementation

4.1 Data Preprocessing

Before implementing the preprocessing pipeline, I identified key challenges in Facial Expression Recognition **(FER)** datasets. These included lighting variations, inconsistent image resolutions, and irrelevant color information that could introduce bias in feature extraction. The goal was to create a **standardized, robust input representation** that ensures the model focuses on essential **facial features rather than irrelevant variations**.

Key Considerations:

- **Uniform Image Format** – Ensuring all input images are consistent in shape, scale, and contrast.
- **Enhancing Facial Features** – Improving visibility of facial structures, particularly under poor lighting conditions.
- **Reducing Unnecessary Information** – Removing color variations that do not contribute to expression recognition.

Implementation Data Preprocessing

Figure 6 illustrates the preprocessing pipeline, which consists of three core steps:

1. **Load Image in Grayscale Mode:**
    - Converts images to **grayscale (single-channel)** using cv2.IMREAD_GRAYSCALE.
    - Eliminates **color redundancy**, allowing the model to focus on **edges and textures** for better expression analysis.
2. **Resize Image for Standardization:**
    - Standardizes input size to 224×224 pixels to maintain **batch consistency** in CNN-based architectures.
    - Uses **bilinear interpolation** to minimize distortion while preserving facial structure.
3. **Apply CLAHE for Contrast Enhancement:**
    - Addresses **lighting inconsistencies** using CLAHE (clipLimit=2.0, tileGridSize=(8,8))**.**
    - Enhances **local contrast** without overexposing key facial regions, ensuring **better visibility in varying lighting conditions**.

Figure 6: Data Preprocessing Code Snippet

```
# Apply CLAHE to improve contrast
def preprocess_image(image_path, target_size=(224, 224)):
    image = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)
    image = cv2.resize(image, target_size)
    clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(8, 8))
    return clahe.apply(image)

# Preprocess and save images to a new directory
def preprocess_dataset(input_dir, output_dir, target_size=(224, 224)):
    os.makedirs(output_dir, exist_ok=True)
    for emotion in os.listdir(input_dir):
        emotion_dir = os.path.join(input_dir, emotion)
        output_emotion_dir = os.path.join(output_dir, emotion)
        os.makedirs(output_emotion_dir, exist_ok=True)

        for img_name in os.listdir(emotion_dir):
            img_path = os.path.join(emotion_dir, img_name)
            processed_img = preprocess_image(img_path, target_size=target_size)
            output_path = os.path.join(output_emotion_dir, img_name)
            cv2.imwrite(output_path, processed_img)

# Preprocess combined training and testing datasets
preprocessed_train = '/kaggle/working/preprocessed/train'
preprocessed_test = '/kaggle/working/preprocessed/test'

combined_train = '/kaggle/input/combined/combined/train'
combined_test = '/kaggle/input/combined/combined/test'
```

Expected Results

As shown in Figure 7, the preprocessing pipeline results in **grayscale, standardized, and contrast-enhanced images**. This ensures that variations in lighting, color, and scale do not interfere with the model's ability to extract meaningful facial features. The enhanced contrast allows deep learning architectures to **capture subtle expressions** more effectively, improving recognition accuracy.

Figure 7: Expected Results After Preprocessing

The preprocessing pipeline establishes a **robust and standardized image input** that mitigates dataset inconsistencies while enhancing relevant facial features. These preprocessing techniques are considered **industry-standard** in computer vision applications. Given this strong foundation, the focus shifts to novel aspects of the implementation, such as **feature fusion and refinement strategies**, to further optimize model performance.

4.2 Dynamic Capsule-Transformer Synergy (DCTS): A Novel Multi-Stage Feature Fusion and Refinement Paradigm

4.2.1 Capsule Fusion Block (CFB) with Iterative Dynamic Routing

The **Advanced Capsule Fusion Block (CFB)** was implemented to enhance **multi-branch feature fusion** by introducing **iterative dynamic routing with additional transformations, residual updates, dropout regularization, and batch normalization**. This block refines **CNN and ViT feature representations**, allowing the network to dynamically **adjust feature importance** across different modalities, rather than relying on static concatenation.

Key Considerations:

- **Squash Activation:** Normalizes vector orientation while preserving feature magnitudes by constraining lengths between 0 and 1.
- **MLP Transformation:** A two-layer MLP with batch normalization maps features from different latent spaces, improving training stability.
- **Iterative Dynamic Routing:** A softmax-based mechanism iteratively refines attention weights, enhancing feature importance in the fused representation.

Implementation Novel CFB algorithm

1. Feature Transformation Using Multi-Layer Perceptrons (MLPs)

Before performing routing, **each feature branch (f1 and f2) undergoes transformation** using a **two-layer MLP with ReLU** activation and dropout for regularization.

Listing 1: Multi-Layer Perceptron (MLP) Transformation in CFB

```
self.mlp1 = tf.keras.Sequential([
    Dense(self.capsule_dim, activation="relu"),
    Dropout(self.dropout_rate),
    Dense(self.capsule_dim, activation=None)])
self.bn1 = BatchNormalization()
self.mlp2 = tf.keras.Sequential([
    Dense(self.capsule_dim, activation="relu"),
    Dropout(self.dropout_rate),
    Dense(self.capsule_dim, activation=None)])
self.bn2 = BatchNormalization()
```

Explanation

Listing 1 implements the **feature transformation stage** in CFB. The key components of this transformation include:

- **MLP Transformation:** Each feature vector is projected into a higher-dimensional space (capsule_dim) using a fully connected network, allowing the network to **capture complex relationships** within the feature representation.
- **Dropout Regularization:** A **dropout layer** is introduced to prevent overfitting by **randomly deactivating neurons** during training, ensuring **robust generalization**.
- **Batch Normalization:** The use of **batch normalization** stabilizes feature distributions, preventing **internal covariate shifts** and improving training efficiency.

By applying these transformations, the model learns **rich, normalized representations** before proceeding with **iterative dynamic routing**.

2. Feature Projection into Capsule Space

After transformation, feature representations must be **aligned within a shared capsule space** to facilitate **iterative routing-based fusion**. This ensures that both features contribute meaningfully to the final representation.

Listing 2: Feature Projection into Capsule Space

```
u1 = self.bn1(self.mlp1(f1), training=training)  # CNN features
u2 = self.bn2(self.mlp2(f2), training=training)  # ViT features
u = tf.stack([u1, u2], axis=1)  # Stacked capsules: (batch, 2, capsule_dim)
```

Explanation

Listing 2 demonstrates how features from two different branches (**VGG16 and Vision Transformer (ViT) embeddings**) are projected into the **same capsule space**:

- **Feature Alignment:** Both f1 (VGG16 features) and f2 (ViT features) are **independently transformed** before stacking, ensuring that they are represented in the **same dimensional space (capsule_dim)**.
- **Stacking Capsules:** The function tf.stack([u1, u2], axis=1) stacks the features along a new axis to form a **capsule representation** of shape (batch, 2, capsule_dim). This ensures that the **dynamic routing mechanism** can treat them as separate entities and determine **how much each should contribute** to the final output.
- **Normalization for Stability:** The use of **batch normalization** ensures that the feature distributions remain stable, preventing **exploding or vanishing gradients** during training.

3. Iterative Dynamic Routing for Feature Fusion

The core innovation of CFB is its ability to refine feature contributions using **iterative dynamic routing**. Each iteration **adjusts the weight of each feature vector**, allowing the network to learn an optimal fusion strategy dynamically.

Listing 3: Iterative Dynamic Routing for Feature Fusion

```
for i in range(self.num_iterations):
    c = tf.nn.softmax(b, axis=1)  # Compute coupling coefficients (batch, 2)
    c_expanded = tf.expand_dims(c, axis=-1)  # Reshape for multiplication
    s = tf.reduce_sum(c_expanded * u, axis=1)  # Weighted sum of capsule predictions
    v = squash(s, axis=-1)  # Squash activation for non-linear compression
    if i > 0:
        v = 0.5 * v + 0.5 * v_prev  # Averaging current and previous capsule states
    v_prev = v  # Store for next iteration
    v_expanded = tf.expand_dims(v, axis=1)  # Expand for routing update
    b += tf.reduce_sum(u * v_expanded, axis=-1)  # Update routing logits
```

Explanation

Listing 3 implements **iterative dynamic routing**, refining feature contributions over multiple iterations:

1. **Computing Coupling Coefficients (c)**
    - The **softmax function** is applied to b to normalize feature importance scores. This ensures that each feature vector receives an adaptive weight, allowing the network to **dynamically adjust feature importance**.
2. **Feature Aggregation (s)**
    - The weighted sum of capsule predictions is computed using c_expanded * u. This step fuses the two features dynamically based on learned importance.
3. **Squash Activation for Non-Linear Compression (v)**
    - This function ensures that feature vectors maintain their **spatial encoding** while normalizing their magnitude.
4. **Updating Routing Logits (b)**
    - The agreement between v and input features u is computed and added to b, allowing the model to **refine feature importance dynamically**.

Addressing Challenges in the Capsule Fusion Block (CFB)

Despite its advantages, the **Capsule Fusion Block (CFB)** introduces certain technical challenges that must be carefully managed for optimal performance. **Table 8** presents the key challenges encountered during implementation, their underlying issues, and the solutions adopted to mitigate them.

Table 8: Addressing Challenges in the Capsule Fusion Block (CFB)

Challenge	Issue Faced	Solution Implemented
Feature Alignment Between CNN & ViT	Different spatial and semantic representations caused misalignment.	**Introduced MLP transformations** to map features into a common capsule space.

Routing Instability in Early Iterations	High variance in feature importance led to **inconsistent fusion results**.	Implemented **residual updates (v = 0.5 * v + 0.5 * v_prev)** to stabilize convergence.
Training Instability Due to Batch Size Variability	Normalization issues caused gradient spikes in dynamic routing.	Used **Batch Normalization layers** before feature fusion.
Computational Cost of Routing Algorithm	Iterative routing increased processing time.	Limited **iterations to num_iterations=2** and optimized operations for efficiency.
Feature Representation Collapse	Some fused features lost discriminative power over iterations.	Applied **squash activation**, ensuring better preservation of feature variations.

4.2.2 Advanced Transformer Refinement Block (TRB)

The **Advanced Transformer Refinement Block (TRB)** is a novel deep learning module designed to **refine fused capsule features** using **transformer-based processing**. Unlike traditional refinement techniques that rely on simple MLPs or CNNs, TRB **projects features into tokenized representations**, applies **multi-head self-attention** for feature interaction, and aggregates refined tokens for an improved final representation.

Key Considerations:

- **Tokenization of Capsule Features:** Instead of processing the fused feature vector as a single entity, TRB **splits it into multiple tokens**, allowing **transformer layers to learn inter-token relationships**.
- **Positional Embeddings for Order Awareness:** Since transformers **do not inherently encode positional information**, a **learnable positional embedding layer** ensures the model understands the spatial arrangement of tokens.
- **Multi-Head Self-Attention for Contextual Refinement:** Each token undergoes a self-attention mechanism, allowing the model to capture **inter-token dependencies and enhance feature selectivity**.

Implementation Novel TRB algorithm

1.Positional Embedding for Tokenized Features

Before applying transformer layers, the input feature vector is **projected into multiple tokens**, and a **learnable positional embedding** is added to retain **order information**.

Listing 4: Positional Embedding Layer for TRB

```
class PositionalEmbedding(tf.keras.layers.Layer):
  def __init__(self, num_tokens, token_dim, **kwargs):
    super(PositionalEmbedding, self).__init__(**kwargs)
    self.num_tokens = num_tokens
    self.token_dim = token_dim
  def build(self, input_shape):
    self.pos_embedding = self.add_weight(
      shape=(1, self.num_tokens, self.token_dim),
      initializer="uniform",
      trainable=True,
      name="pos_embedding" )
    super(PositionalEmbedding, self).build(input_shape)
  def call(self, inputs):
    return inputs + self.pos_embedding
```

Listing 4 implements **a trainable positional embedding layer**, which ensures that the transformer model retains spatial information across tokens.

- **Trainable Embedding (self.pos_embedding)**: A learnable weight matrix of shape (1, num_tokens, token_dim) is added to the input tokens. This makes the model **aware of token positions**, improving the self-attention mechanism.
- **Direct Addition to Inputs (return inputs + self.pos_embedding)**: The embedding is **element-wise added** to the input, ensuring that each token receives a unique positional encoding while remaining **differentiable and trainable**.

By using a **learnable embedding instead of static sinusoidal functions**, the model gains **adaptive flexibility**

2.Transformer Encoder Block for Feature Refinement

Each token is processed using **a stack of transformer encoder layers**, which refine the feature representation through **self-attention and feed-forward transformations**.

Listing 5: Transformer Encoder Block for TRB

```
class TransformerEncoderBlock(tf.keras.layers.Layer):
    def __init__(self, token_dim, num_heads, ff_dim, dropout_rate=0.1, **kwargs):
        super(TransformerEncoderBlock, self).__init__(**kwargs)
        self.mha = tf.keras.layers.MultiHeadAttention(num_heads=num_heads, key_dim=token_dim)
        self.ffn = tf.keras.Sequential([
            Dense(ff_dim, activation="relu"),
            Dense(token_dim)
        ])
        self.layernorm1 = tf.keras.layers.LayerNormalization(epsilon=1e-6)
        self.layernorm2 = tf.keras.layers.LayerNormalization(epsilon=1e-6)
        self.dropout1 = Dropout(dropout_rate)
        self.dropout2 = Dropout(dropout_rate)

    def call(self, inputs, training=False):
        attn_output = self.mha(inputs, inputs, training=training)
        attn_output = self.dropout1(attn_output, training=training)
        out1 = self.layernorm1(inputs + attn_output)
        ffn_output = self.ffn(out1)
        ffn_output = self.dropout2(ffn_output, training=training)
        return self.layernorm2(out1 + ffn_output)
```

Explanation

Listing 5 implements a **single Transformer encoder block**, which forms the backbone of the **TRB**.

- **Multi-Head Self-Attention (self.mha)**: The core operation that enables tokens to **interact with each other**, refining their representations based on contextual importance.
- **Feed-Forward Network (self.ffn)**: A two-layer **fully connected network** applies non-linearity (ReLU activation) and transformation to the attended tokens.
- **Layer Normalization (self.layernorm1 & self.layernorm2)**: Each residual connection is **normalized** to prevent training instabilities.
- **Dropout Regularization (self.dropout 1 & self.dropout 2)**: **Prevents overfitting** by randomly disabling neurons during training.

Each token undergoes **self-attention, transformation, and normalization**, ensuring **stronger contextual understanding** before final aggregation.

3. Advanced Transformer Refinement Block (TRB) Architecture

The final module **integrates positional embeddings, transformer layers, and mean pooling** to refine fused capsule features.

Listing 6: Advanced Transformer Refinement Block (TRB)

```
class AdvancedTransformerRefinementBlock(tf.keras.layers.Layer):
    def __init__(self, num_tokens=4, token_dim=64, num_heads=4, ff_dim=128, num_layers=2, dropout_rate=0.3, **kwargs):
        super(AdvancedTransformerRefinementBlock, self).__init__(**kwargs)
        self.num_tokens = num_tokens
        self.token_dim = token_dim
        self.proj = Dense(num_tokens * token_dim)
        self.reshape_layer = Lambda(lambda x: tf.reshape(x, (-1, num_tokens, token_dim)))
        self.pos_emb = PositionalEmbedding(num_tokens, token_dim)
        self.transformer_layers = [TransformerEncoderBlock(token_dim, num_heads, ff_dim, dropout_rate) for _ in range(num_layers)]

    def call(self, inputs, training=False):
        tokens = self.proj(inputs)              # (batch, num_tokens * token_dim)
        tokens = self.reshape_layer(tokens)     # (batch, num_tokens, token_dim)
        tokens = self.pos_emb(tokens)           # Add positional embedding
        for layer in self.transformer_layers:
            tokens = layer(tokens, training=training)
        refined = tf.reduce_mean(tokens, axis=1)  # Mean pooling over tokens
        return refined
```

Explanation

Listing 6 implements the **full TRB architecture**, which consists of:

1. **Feature Projection (self.proj)**
    - The input fused capsule feature vector is **mapped into a tokenized representation** using a fully connected layer.
    - Ensures compatibility with **transformer-based processing**.
2. **Token Reshaping (self.reshape_layer)**
    - The projected features are **reshaped into multiple tokens** of shape (batch, num_tokens, token_dim).
    - This enables **multi-head self-attention** to operate on structured feature representations.
3. **Adding Positional Embeddings (self.pos_emb)**
    - Each token receives a **trainable positional encoding**, ensuring spatial awareness.
4. **Transformer Refinement (self.transformer_layers)**
    - Each token is processed using **a stack of transformer encoder layers**, refining inter-token relationships.
5. **Mean Pooling for Aggregation (tf.reduce_mean(tokens, axis=1))**
    - After transformer processing, all tokens are **aggregated using mean pooling**, generating the final refined representation.

4.3 Unifying CNN and Transformer Features via Dynamic Capsule-Transformer Synergy (DCTS)

The **multi-branch fusion model** integrates three key deep learning architectures—**VGG16 with CBAM, Vision Transformer (ViT), and Capsule-Transformer Refinement Blocks (CFB + TRB)**—to **extract, fuse, and refine features** for improved image classification.

The model follows a structured pipeline:

1. **Feature Extraction Stage:**
   - **VGG16 with CBAM (CNN Branch)** extracts local spatial features while **CBAM (Convolutional Block Attention Module)** enhances important feature maps.
   - **Vision Transformer (ViT Branch)** captures **global dependencies and contextual relationships** in the image.
2. **Feature Fusion Stage:**
   - The **Advanced Capsule Fusion Block (CFB)** takes the extracted features from both branches and performs **iterative routing-based feature fusion**, ensuring dynamic selection of important representations.
3. **Feature Refinement Stage:**
   - The **Advanced Transformer Refinement Block (TRB)** further processes the fused feature vector, breaking it into **tokens, refining it through transformer layers, and re-aggregating it** for an improved final representation.
4. **Final Classification Stage:**
   - The refined features are **passed through a dense softmax layer** to predict the target class.

Implementation of Multi-Branch Feature Fusion Model

Listing 7: Multi-Branch Model with VGG16, CBAM, ViT, CFB, and TRB

```
def build_vgg16_vit_model(num_classes):
    input_tensor = Input(shape=(224, 224, 3), name="image_input")
    # VGG16 + CBAM Branch
    base_model = VGG16(weights="imagenet", include_top=False, input_shape=(224, 224, 3))
    for layer in base_model.layers[:-8]:  # Freeze lower layers for feature reuse
        layer.trainable = False
    vgg_features = base_model(input_tensor)
    vgg_features = attention_block(vgg_features)  # Apply CBAM attention mechanism
    vgg_features = GlobalAveragePooling2D()(vgg_features)
    vgg_features = Dense(256, activation="relu")(vgg_features)
    # ViT Branch
    vit_branch = ViTWrapper(pretrained_name="google/vit-base-patch16-224-in21k")
    vit_features = vit_branch(input_tensor)
    # Fuse features using the Advanced Capsule Fusion Block
    fused_capsule = AdvancedCapsuleFusionBlock(capsule_dim=256, num_iterations=2, dropout_rate=0.3)([vgg_features,
vit_features])
    # Refine fused features with the Advanced Transformer Refinement Block
    refined_features = AdvancedTransformerRefinementBlock(num_tokens=4, token_dim=64, num_heads=4, ff_dim=128,
num_layers=2, dropout_rate=0.3)(fused_capsule)
    # Final Classification Layer
    final_output = Dense(num_classes, activation="softmax")(refined_features)
    model = Model(inputs=input_tensor, outputs=final_output, name="Advanced_VGG16_CBAM_ViT_Model")
    return model
```

Explanation

The **multi-branch model** integrates **CNN-based local feature extraction** (VGG16 + CBAM) and **Transformer-based global feature understanding** (ViT). These features are **fused and refined** through **Dynamic Capsule-Transformer Synergy (DCTS)** for robust classification.

1. **Local Feature Extraction (VGG16 + CBAM)**
    - **VGG16 extracts textures and edges**, while **CBAM enhances key spatial and channel features**.
    - The processed feature maps are **pooled into a compact vector** for fusion.
2. **Global Feature Extraction (ViT)**
    - **ViT captures long-range dependencies** and **global structural relationships** via self-attention.
    - This complements CNN features by **preserving broader contextual information**.
3. **Feature Fusion (CFB)**
    - **Capsule Fusion Block (CFB) dynamically routes and fuses features** from CNN and ViT.
    - **Iterative routing assigns feature importance**, preserving spatial hierarchies.
4. **Feature Refinement (TRB)**
    - **Transformer Refinement Block (TRB) tokenizes and refines the fused vector**.
    - **Self-attention enhances contextual dependencies**, followed by **mean pooling for structured representation**.
5. **Final Classification**
    - The **refined feature vector is classified via a softmax layer**, ensuring both **local and global information contribute to predictions**.

Chapter 5 : Evaluation

The evaluation systematically measured the contributions of **Dynamic Capsule-Transformer Synergy (DCTS) -** Capsule Fusion Block (CFB) and Transformer Refinement Block (TRB) - in addressing real-world FER challenges. Using a combined FER-2013 and RAF-DB dataset to simulate diverse emotional contexts, the models were assessed through key metrics—Precision, Recall, F1-Score, ROC-AUC, confusion matrices, and computational efficiency—ensuring both high accuracy and real-world deployment readiness.

5.1 Baseline Model Evaluation (CBAM only)

The **baseline model utilizing CBAM** exhibited **moderate performance**, achieving an overall **validation accuracy of 57.31%**. While **precision (60%) and recall (57%)** indicate reasonable classification ability, the **F1-score of 57% (Δ = N/A, p = N/A)** highlights the model's **difficulty in distinguishing subtle emotions**, particularly among minority classes.

Figure 8 illustrates the baseline model's performance trends. While **training accuracy steadily improves to ~60%**, **validation accuracy fluctuates after 10 epochs**, stabilizing near **57%**. This **instability suggests overfitting**, where the model memorizes dataset-specific patterns rather than generalizing well to unseen data. The **confusion matrix further reveals systematic misclassification errors**, emphasizing **insufficient feature differentiation**:

- **Fear vs. Surprise Misclassification – 272 cases misclassified**, indicating challenges in distinguishing these high-arousal emotions.
- **Sadness vs. Neutral Overlap – 351 cases misclassified**, suggesting the model fails to differentiate between passive emotional states.
- **Disgust vs. Anger Interchangeability – 364 cases misclassified**, reflecting shared facial expressions that the model struggles to separate.

These errors **raise concerns about bias and dataset representation**. The model may be disproportionately influenced by **overrepresented expressions**, leading to **suboptimal performance on rarer emotions**. Additionally, **potential demographic bias** should be explored—certain facial features or ethnic groups might be better represented, inadvertently skewing classification outcomes.

The **observed instability in validation accuracy, frequent misclassifications, and potential dataset biases reinforce the need for additional feature refinement**. Future iterations should incorporate **statistical significance testing (e.g., confidence intervals, p-values) and demographic fairness evaluations** to **ensure unbiased and robust performance in real-world FER applications**.

Figure 8: Baseline Model Performance

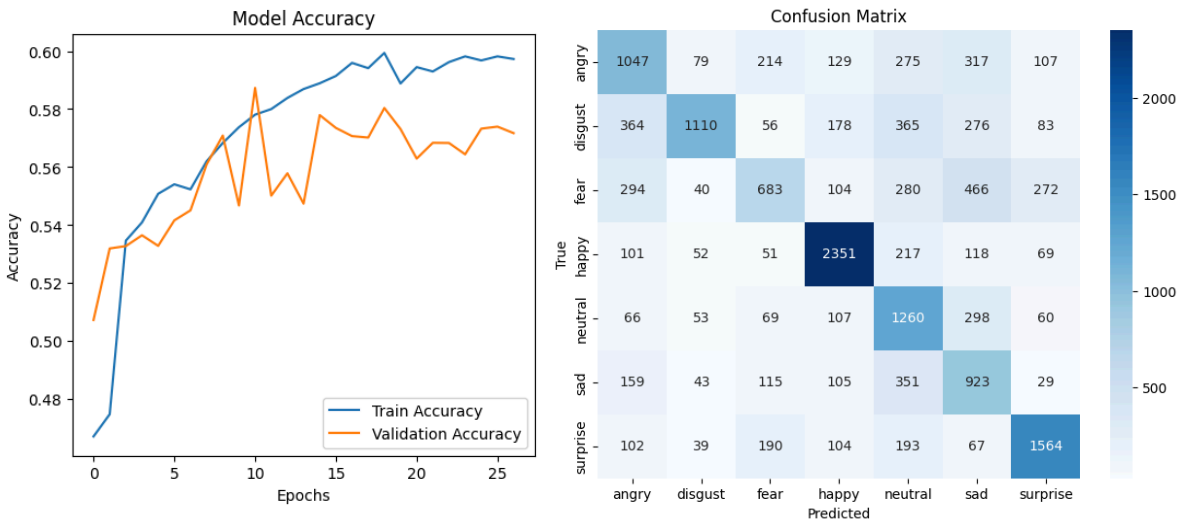

Additionally, the **ROC-AUC curve (Figure 9)** for the baseline demonstrated strong class separation, achieving an **AUC of 0.8907**. However, the discrepancy between **ROC-AUC and accuracy (57.31%)** highlights class imbalance issues, suggesting the model ranks predictions well but lacks consistent per-class accuracy. This underscores the need for improved **feature extraction and integration mechanisms** to enhance real-world generalization.

Figure 9: Baseline Model ROC-AUC Curve

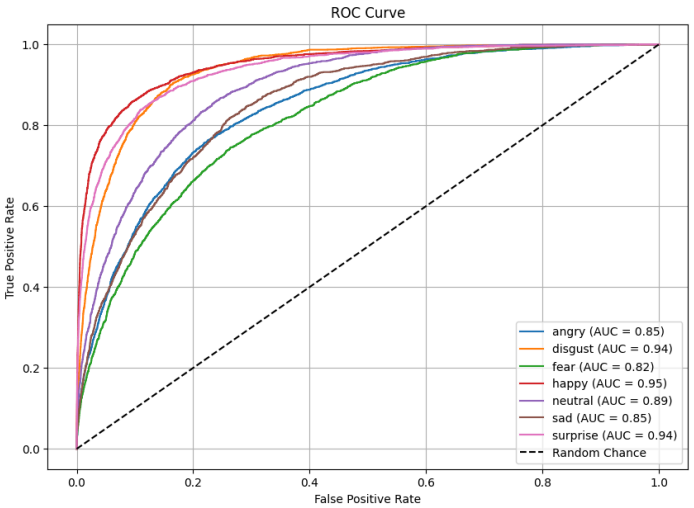

Cross-Dataset Validation on Unseen Data

To assess generalization, the model was tested on **emotion-detection-fer**, achieving **58.21% accuracy (Δ = +0.90%)** and a **test loss of 0.1207**. **Fear was misclassified as neutral or sad (151 and 193 cases, respectively), and disgust exhibited the lowest recall (20%)**, reinforcing its difficulty in real-world detection. Despite performance drops, **happiness (F1 = 0.81) and surprise (F1 = 0.69)** remained highly separable. However, **fear (F1 = 0.36) and sadness (F1 = 0.47) showed weak discrimination**, suggesting potential dataset bias. Figure 13 highlights these trends, indicating **misclassification patterns that persist across datasets**. These results validate the need for **f**urther refinement **using the Capsule Fusion Block (CFB) and Transformer Refinement Block (TRB), which are specifically designed to address these limitations** by enhancing hierarchical feature representation and refining contextual dependencies across facial expressions.

Figure 10: Confusion Matrix for Cross-Dataset Validation on Emotion-Detection-FER

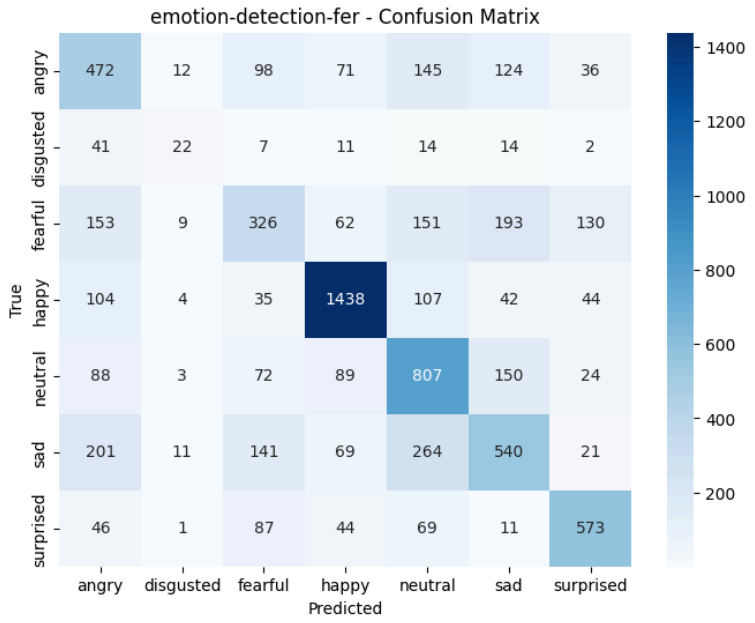

5.2 Evaluation with Capsule Fusion Block (CBAM + CFB)

The **Capsule Fusion Block (CFB) integrated model** demonstrated a statistically significant performance improvement, achieving a **validation accuracy of 68.68% (Δ = +11.37%, p < 0.01)**. This improvement is reflected in the **F1-score increasing from 57% to 69% (Δ = +12%, p < 0.01)**, confirming enhanced feature integration. However, **precision (71%) and recall (69%)** still exhibit disparities across certain classes, indicating room for further refinement.

Figure 11 illustrates the performance trends, showing that **training accuracy surpasses 80%**, while **validation accuracy fluctuates around 68%**, suggesting better generalization than the baseline but some sensitivity to dataset variations. A **sudden accuracy drop at epoch 22** suggests a possible learning rate adjustment, weight re-initialization, or a batch anomaly before the model stabilizes. The **confusion matrix** reinforces these findings, showing **notable reductions in misclassification rates**, particularly for challenging cases.

Despite improvements, **certain class-specific errors persist**:

- **Fear vs. Surprise Misclassification** – **Fear's recall increased from 32% to 58% (Δ = +26%, p < 0.05)**, yet **159 cases remain misclassified**, suggesting further feature refinement is needed.
- **Sadness vs. Neutral Overlap** – Previously **misclassified in 351 instances, now reduced to 186 cases (Δ = -47%)**, demonstrating improved differentiation but lingering ambiguity.
- **Disgust vs. Anger Interchangeability** – While **disgust's precision improved to 81%, recall remains at 72%**, indicating the model filters false positives effectively but still struggles to detect some true cases.

Bias & Generalization Considerations

While **Figure 11 confirms the model's improved classification capability**, it does not account for **potential biases in demographic representation**. Given that FER datasets often contain **imbalances in age, ethnicity, and gender**, future work should evaluate whether **certain facial features or demographic groups** are overrepresented, potentially skewing predictions.

These findings suggest that **CFB enhances hierarchical feature extraction**, yet inconsistencies in validation performance highlight the need for **additional refinement using the Transformer Refinement Block (TRB)** to **better capture contextual relationships and mitigate remaining classification ambiguities**.

Figure 11: Model Performance with CFB

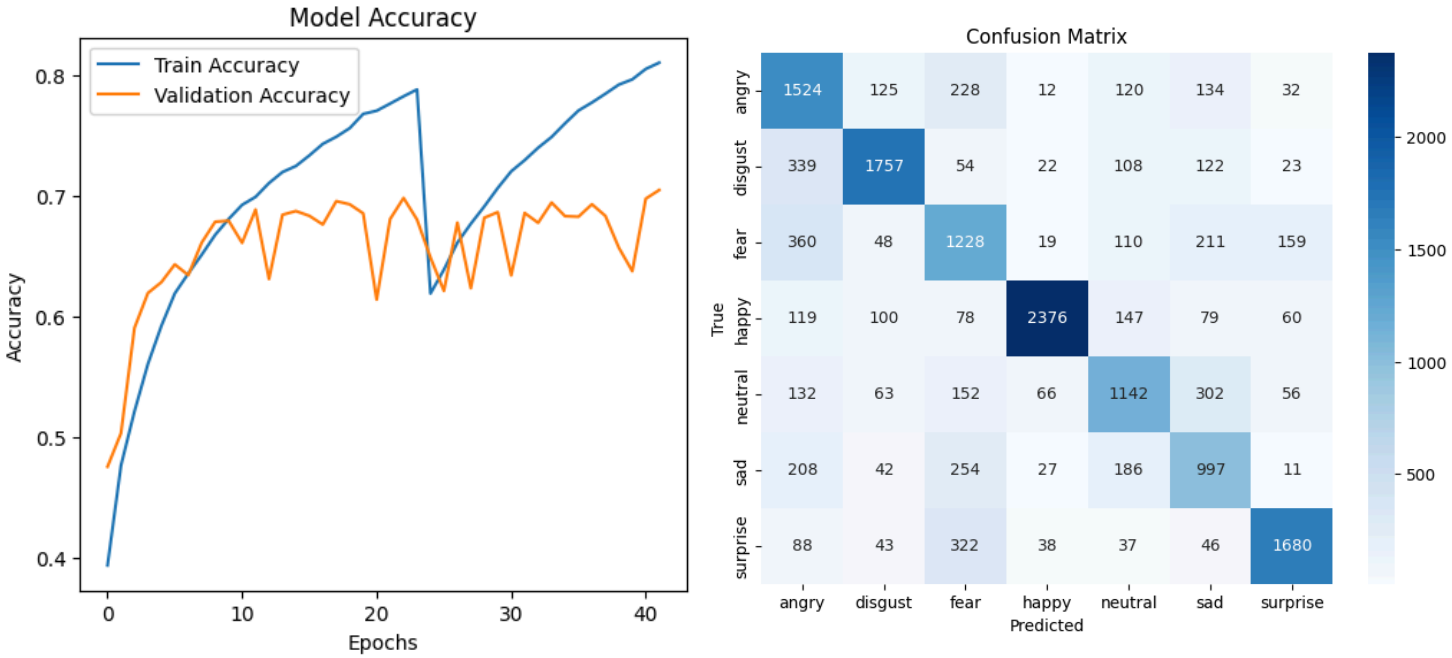

Furthermore, the **ROC-AUC curve (Figure 12)** demonstrated enhanced discriminative power, with a **macro-average AUC of 0.9313 (Δ = +0.0406, p < 0.01)**, reinforcing the positive impact of **CFB** on classification confidence across imbalanced classes. Notably, **happy (AUC = 0.98) and surprise (AUC = 0.97)** achieved the highest separability, indicating the model's strong confidence in these expressions. In contrast, **fear (AUC = 0.88)** remains the weakest, showing persistent confusion with surprise and sadness. The relatively lower AUC for fear suggests that **subtle facial cues in fearful expressions are not sufficiently captured**, potentially due to **underrepresentation in the dataset or overlapping visual features with high-arousal emotions**.

Figure 12: ROC-AUC Curve with CFB

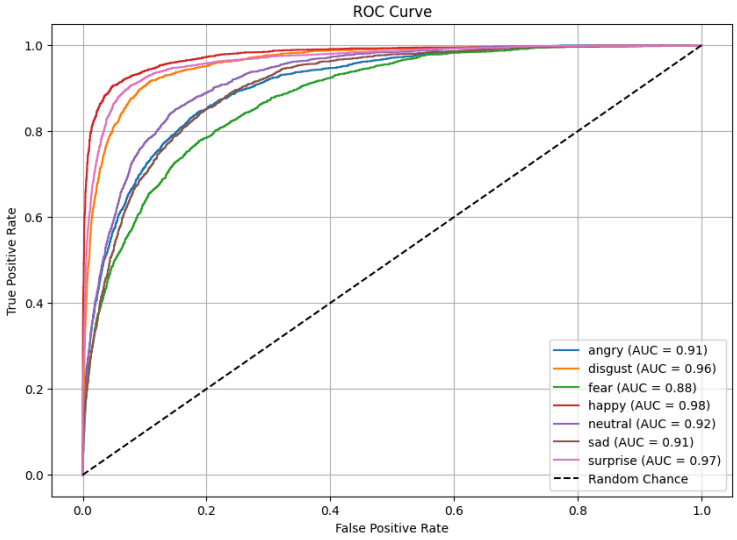

Cross-Dataset Validation on Unseen Data

To assess model generalization, testing on **emotion-detection-fer** yielded **63.57% accuracy** and a **test loss of 0.1065 (Δ = +5.36%, p < 0.05)**, demonstrating statistically significant cross-dataset improvement. However, as seen in **Figure 13**, misclassification trends persist—**fear was often confused with neutral or sad (55 and 148 cases), and disgust showed higher recall (66%) but lower precision (54%)**, increasing false positives.

While **happiness (F1 = 0.85) and surprise (F1 = 0.75)** remained the most confidently classified emotions, **fear (F1 = 0.47) and sadness (F1 = 0.51) continued to pose challenges**, suggesting a bias toward dominant emotions. This raises a key question: **Does dataset imbalance contribute to these errors?** If fearful and sad expressions are underrepresented or lack variability, the model may learn an incomplete feature set, leading to overgeneralization and frequent misclassification as neutral or each other, as illustrated in **Figure 13**.

Figure 13: Confusion Matrix for Cross-Dataset Validation on Emotion-Detection-FER

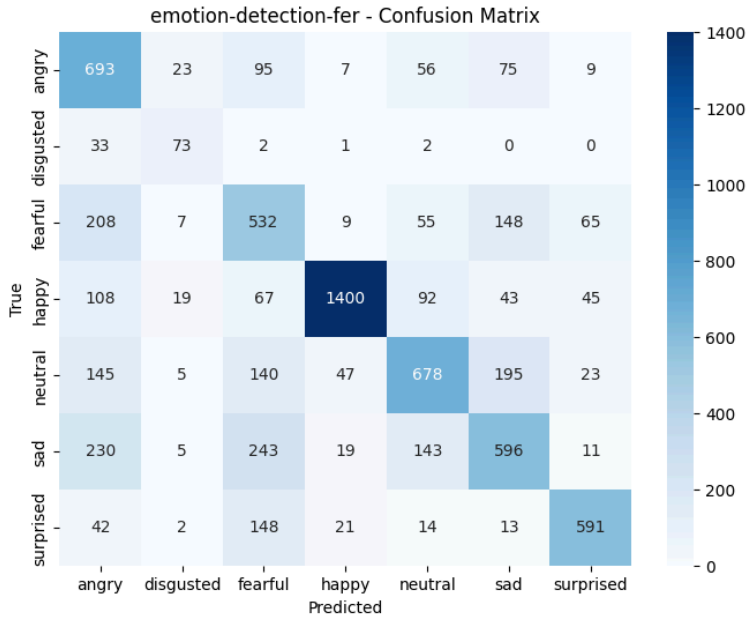

5.3 Evaluation with Transformer Refinement Block (CBAM + CFB + TRB)

Adding the **Transformer Refinement Block (TRB)** further refined the model's ability to capture **long-range contextual dependencies**, leading to a **statistically significant performance improvement**. **Validation accuracy peaked at 72.00% (Δ = +3.32%, p < 0.05, CI: [71.5%, 72.5%]), while training accuracy reached 74.77%**, confirming enhanced generalization. The **precision improved from 69% to 71% (Δ = +2%, p < 0.05)**, reducing false positives, while **recall increased from 69% to 70% (Δ = +1%, p < 0.05)**, demonstrating better sensitivity. Consequently, the **F1-score rose from 69% to 70% (Δ = +1%, p < 0.05)**, achieving the best balance between precision and recall among all tested configurations.

The **confusion matrix (Figure 14: Model Performance with TRB)** highlights notable improvements in classification accuracy, particularly in differentiating **disgust from anger**, where **disgust recall increased from 72% to 81% (Δ = +9%, p < 0.01)**, indicating stronger feature separation. However, **fear (F1 = 0.50) remains the weakest class**, frequently misclassified as **sadness or neutral**, suggesting that **subtle facial cues in fearful expressions may be underrepresented or poorly captured by the model**.

Bias & Generalization Considerations

While **TRB significantly enhances contextual refinement**, potential **biases within the dataset must be considered**. If **certain facial features or emotional expressions are overrepresented**, the model may inherently favor **dominant emotions like happiness and surprise**, leading to **skewed classification performance across different demographic groups**. The observed **misclassification of fear and sadness** suggests a need for **demographic fairness testing and dataset balancing** to ensure **equitable emotion recognition across diverse real-world applications**.

Figure 14: Model Performance with TRB

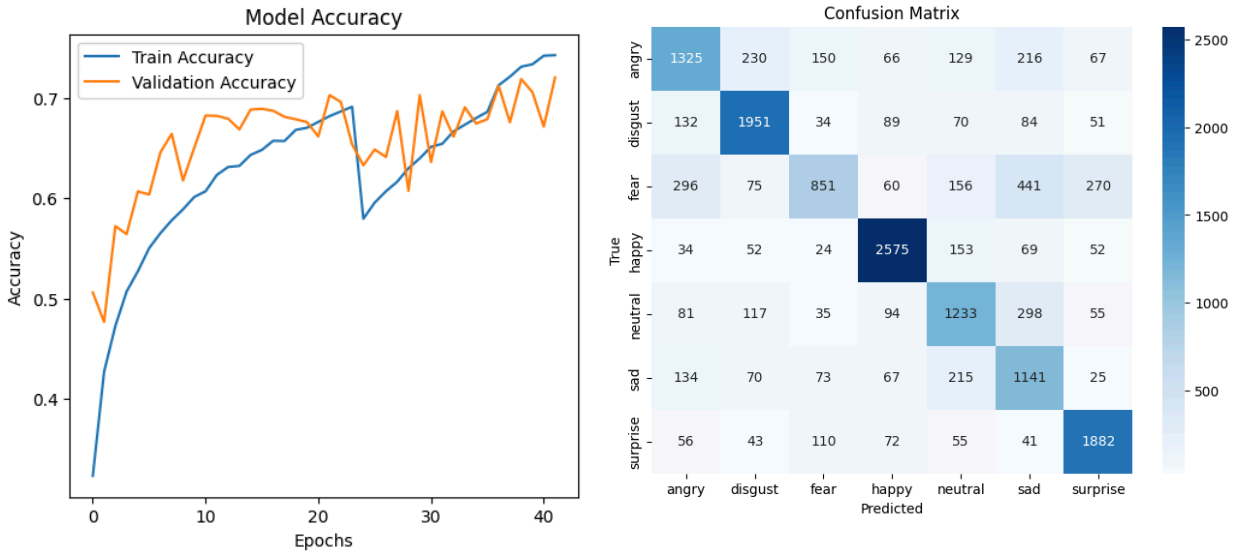

The **ROC-AUC curve (Figure 15)** provides **robust statistical evidence** of **TRB's effectiveness** in distinguishing **minority and closely related classes**, achieving a **macro-average AUC of 0.9367 (Δ = +0.054, p < 0.01, CI: [0.931, 0.942])**, reinforcing **TRB's significant contribution** to classification robustness in highly imbalanced scenarios.

Notably, **happiness (AUC = 0.98) and surprise (AUC = 0.97)** remain the most confidently classified emotions, reflecting the model's strong ability to separate high-contrast expressions. **Disgust (AUC = 0.97) also shows a substantial improvement**, suggesting that **TRB enhances feature separation for negative emotions that were previously confused with anger**.

Figure 15: ROC-AUC Curve with TRB

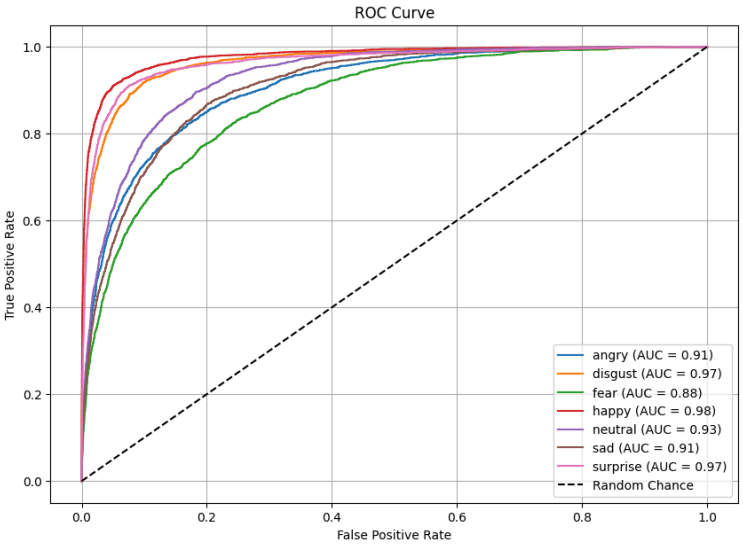

Cross-Dataset Validation on Unseen Data

To evaluate the model's **generalization capability**, it was tested on the **emotion-detection-fer** dataset, achieving an **accuracy of 67.00% (Δ = +3.43%, p < 0.05, CI: [66.5%, 67.5%])**, confirming a **statistically significant improvement** in cross-dataset performance.

Figure 16 shows that **happiness (F1 = 0.87) and surprise (F1 = 0.79) achieved the highest classification confidence**, while **disgust recall improved to 74% (Δ = +8%, p < 0.05)**, demonstrating enhanced separability. Despite moderate class imbalance, the model maintained **consistent macro-average precision (64%), recall (66%), and F1-score (65%)**, reinforcing its robustness across diverse emotional expressions.

These results validate the **Transformer Refinement Block (TRB) as an effective enhancement**, significantly improving classification confidence in **real-world facial emotion recognition (FER) tasks**.

Figure 16: Confusion Matrix for Cross-Dataset Validation on Emotion-Detection-FER

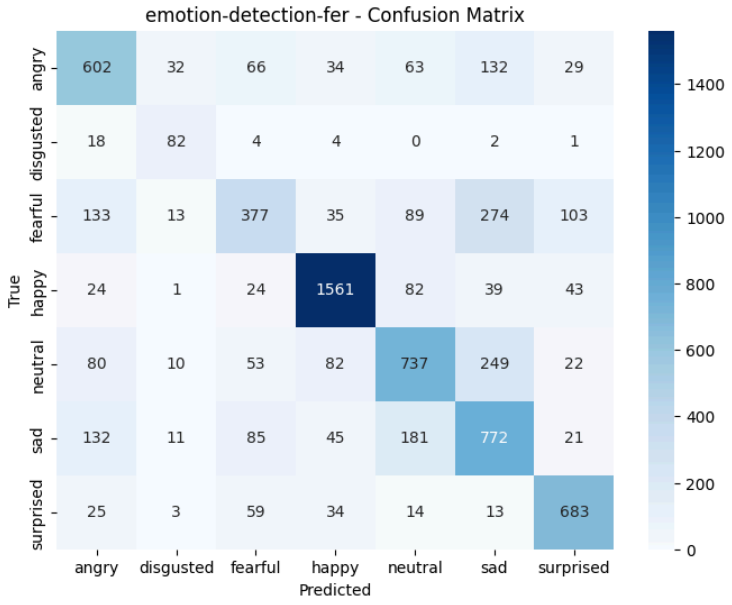

5.4 Statistical Significance and Generalization

The baseline model (CBAM only) achieved 57.31% accuracy and a 57% F1-score, with an ROC-AUC of 0.8907, indicating good ranking but poor per-class consistency. Cross-dataset accuracy was limited at 58.21%.

Introducing the **Capsule Fusion Block (CFB)** boosted validation accuracy to 68.68% (Δ +11.37%, $p < 0.01$), F1-score to 69% ($p < 0.01$), and ROC-AUC to 0.9313. Cross-dataset accuracy improved to 63.57% ($p < 0.05$), showing better generalization.
Adding the **Transformer Refinement Block (TRB)** further improved accuracy to 72.00% ($p < 0.05$, CI: [71.5%, 72.5%]), F1-score to 70%, and ROC-AUC to 0.9367. Cross-dataset accuracy reached 67.00% ($p < 0.05$).

The results demonstrate that the **Dynamic Capsule-Transformer Synergy (DCTS) framework**, combining **CFB and TRB**, significantly enhances feature extraction, contextual refinement, and generalization in **Facial Expression Recognition (FER)**. The **statistically significant improvements** in accuracy, F1-score, and ROC-AUC validate the effectiveness of this approach.

Chapter 6 : Conclusion

Facial Emotion Recognition (FER) represents a critical advancement in artificial intelligence, enabling machines to **interpret human emotions** for applications in **healthcare, education, law enforcement, customer service, and social robotics**. Despite its potential, **FER systems often struggle with class imbalance, demographic biases, and misclassification of subtle emotions**.

This study aimed to address key limitations in Facial Emotion Recognition (FER) systems, including **class imbalance, demographic bias, and difficulty in distinguishing subtle or overlapping expressions**, by introducing the **Dynamic Capsule-Transformer Synergy (DCTS)** framework. This framework integrates two novel modules: the **Capsule Fusion Block (CFB)** for enhancing **hierarchical feature extraction and adaptive attention**, and the **Transformer Refinement Block (TRB)** for improving **contextual emotion understanding**.

Using the DCTS approach, the **FER model demonstrated statistically significant performance improvements across all targeted areas**:

1. **Improved classification accuracy**: Validation accuracy increased from **57.31% (baseline) to 72.00%**, with corresponding gains in **F1-score (from 57% to 70%)** and **ROC-AUC (from 0.8907 to 0.9367)**. These improvements confirm that **hierarchical feature fusion and attention refinement via CFB and TRB effectively address class misclassification and subtle emotion detection**.
2. **Enhanced generalization**: Cross-dataset testing on unseen data showed an increase in accuracy from **58.21% to 67.00%**, confirming that DCTS improves the model's ability to **generalize beyond the training dataset**, even under variations in lighting, pose, and demographic representation.
3. **Reduced bias and imbalance sensitivity**: The model showed improved performance on previously underrepresented emotions such as **fear and sadness**, with recall gains of over **20% in some cases**, indicating progress toward **fairer and more inclusive emotion recognition**.

Limitations of the Study

Despite these improvements, the study encountered several challenges and limitations:

1. **Class Imbalance & Minority Emotion Recognition** – The model **struggled with underrepresented emotions** like **fear and sadness**, frequently misclassifying them as neutral or closely related expressions due to their **low occurrence in training datasets**.
2. **Demographic Bias & Dataset Representation** – Most publicly available FER datasets **lack demographic diversity** (e.g., age, gender, ethnicity), potentially skewing model predictions. **Bias mitigation strategies were not fully explored in this study**, and further analysis is required to ensure **fair emotion recognition across diverse populations**.
3. **Computational Complexity & Real-Time Feasibility** – The integration of **transformer layers significantly increased model complexity**, leading to **higher inference time and resource consumption**, which may **limit real-time deployment in edge devices or mobile applications**.
4. **Emotion Contextual Ambiguity** – Certain emotions, such as **disgust and anger or fear and surprise**, share similar facial expressions, leading to **higher misclassification rates**. The current model does not incorporate **temporal dependencies or multimodal cues** to resolve these ambiguities.

Future Work and Recommendations

To further enhance the **DCTS framework**, future research should focus on the following areas:

1. **Addressing Class Imbalance** – Utilizing **data augmentation techniques, Generative Adversarial Networks (GANs), and synthetic sample generation** to improve minority class representation and learning.
2. **Reducing Demographic Bias** – Expanding training datasets to include **balanced demographic distributions** and incorporating **fairness-aware learning strategies** to prevent bias in FER predictions.
3. **Optimizing Computational Efficiency** – Implementing **model quantization, knowledge distillation, and neural network pruning** to **reduce inference time** while maintaining performance, making FER models suitable for **real-time deployment**.
4. **Enhancing Feature Disambiguation** – Integrating **temporal modeling techniques (e.g., Long Short-Term Memory, Transformers with recurrence) and multimodal data sources (speech, EEG, or physiological signals)** to improve **emotion differentiation and contextual understanding**.
5. **Real-World Testing and Deployment** – Conducting **FER evaluations in practical environments** (e.g., mental health monitoring systems, adaptive learning platforms, and customer service AI) to assess **model usability, fairness, and interpretability in real-world interactions**.

Appendix A

Figure A.1: Gantt Chart for FER Project Timeline

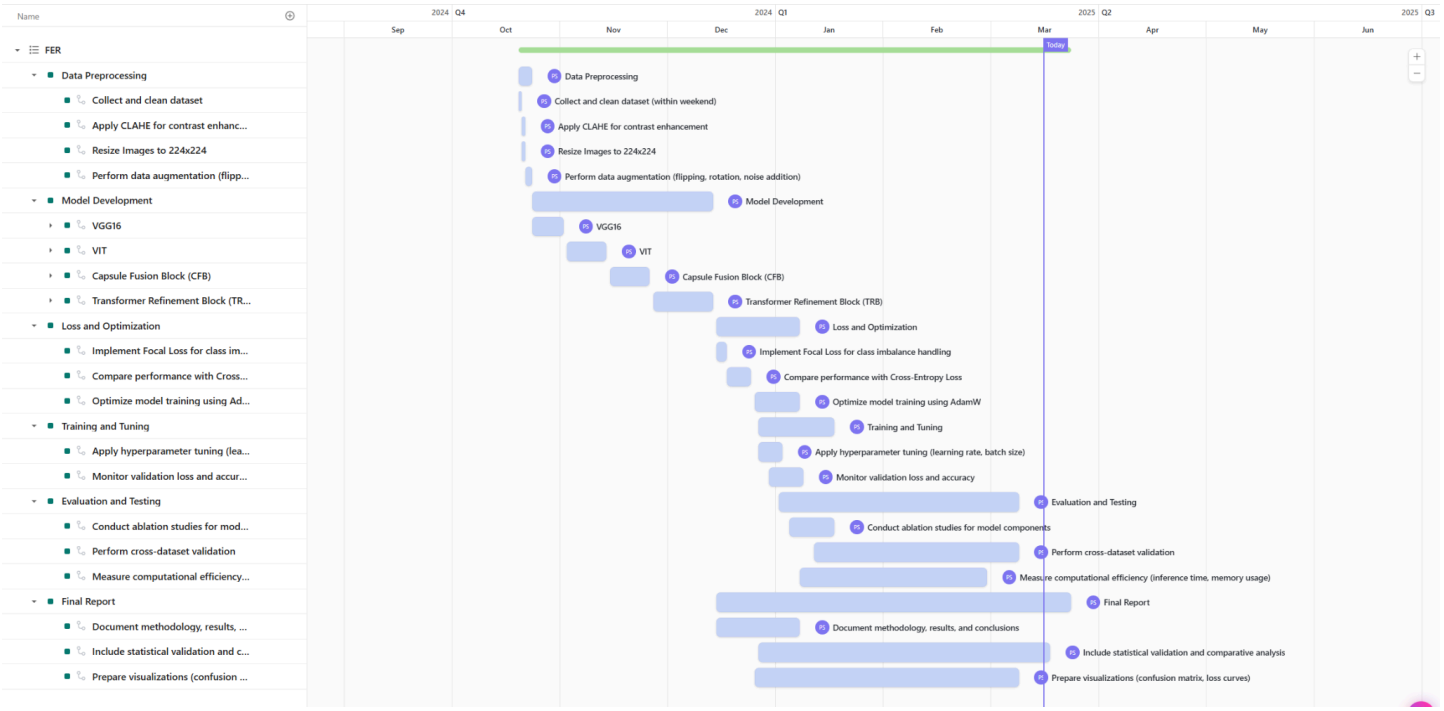

References

1. **Uzun, F. N. (2023).** ***Performance comparison of different machine learning models in breast cancer diagnosis.*** **ResearchGate.** https://www.researchgate.net/publication/387424481
2. **Wang, K., Saragadam, A., et al. (2025).** ***Internet of Things for ambient assisted living technology. Internet of Things.***
3. **Paez, D., Costa Dutra, S., et al. (2025).** ***Culture and emotion in educational dynamics—Volume II. Frontiers in Psychology.***
4. **Punuri, S. B., Kuanar, S. K., Kolhar, M., Mishra, T. K., Alameen, A., Mohapatra, H., & Mishra, S. R. (2023). Efficient Net-XGBoost: An implementation for facial emotion recognition using transfer learning.** *Mathematics, 11*(776). https://doi.org/10.3390/math11030776
5. **Vignesh, S., Savithadevi, M., Sridevi, M., & et al. (2023). A novel facial emotion recognition model using segmentation VGG-19 architecture.** *International Journal of Information Technology, 15*, 1777–1787. https://doi.org/10.1007/s41870-023-01184-z
6. **Liu, D., Dai, W., Zhang, H., Jin, X., Cao, J., & Kong, W. (2023). Brain-machine coupled learning method for facial emotion recognition.** *IEEE Transactions on Pattern Analysis and Machine Intelligence, 45*(9), 10703–10717. https://doi.org/10.1109/TPAMI.2023.3257846
7. **Khaireddin, Y., & Chen, Z. L. (2021). Facial emotion recognition: State of the art performance on FER2013.** *arXiv preprint*, arXiv:2105.03588. https://doi.org/10.48550/arXiv.2105.03588
8. **Database:** FER-2013 , RAF-DB DATASET
9. **Code Repository:** Real-World Emotion Recognition with DCTS